

## **Problem Statement**

In this data science project, the objective is to predict product demand using machine learning. To achieve this goal, we will analyze the provided dataset containing information about product sales, including Store ID, Total Price, Base Price, and Units Sold. The primary challenge is to build a predictive model that can forecast product demand based on the available features.

## **Dataset Details**

### **Dataset Information**

ID: Unique identifier for each sales record.

Store ID: Identifies the store where the product was sold.

Total Price: The actual total price of the product.

Base Price: The base price of the product without any discounts.

Units Sold: The number of units of the product sold in each transaction.

## **Data Exploration**

Before proceeding with the analysis and modeling, we will explore the dataset to gain a better understanding of its characteristics. This includes checking for missing values, data distribution, and statistical summary.

## **Data Preprocessing**

To prepare the data for machine learning, we will need to perform various preprocessing tasks such as handling missing values, encoding categorical features (if any), and scaling/normalizing numerical features. This step is crucial to ensure the data is ready for model training and evaluation.

## **Feature Selection**

Selecting the most relevant features for the prediction task is essential for model accuracy. We will identify and choose the features that have the most impact on product demand.

## Model Building

In this phase, we will build and evaluate machine learning models for product demand prediction. Various regression algorithms will be considered, and their performance will be assessed using appropriate metrics (e.g., Mean Absolute Error, Root Mean Squared Error).

## Model Evaluation

To determine the best model for predicting product demand, we will evaluate the models' performance on a test dataset and compare their results. Model evaluation will involve assessing accuracy, precision, recall, and other relevant metrics.

## Importing Libraries

```
import matplotlib.pyplot as plt
import pandas as pd
import pylab as pl
import numpy as np
import sklearn
%matplotlib inline
```

## Loading the Data Set

```
df = pd.read_csv("ProductDemand.csv")
print(df)
```

	ID	Store ID	Total Price	Base Price	Units Sold
0	1	8091	99.0375	111.8625	20
1	2	8091	99.0375	99.0375	28
2	3	8091	133.9500	133.9500	19
3	4	8091	133.9500	133.9500	44
4	5	8091	141.0750	141.0750	52
...	...	...	...	...	...
150145	212638	9984	235.8375	235.8375	38
150146	212639	9984	235.8375	235.8375	30
150147	212642	9984	357.6750	483.7875	31

150148	212643	9984	141.7875	191.6625	12
150149	212644	9984	234.4125	234.4125	15

[150150 rows x 5 columns]

## Exploring the Data Set

**print(df.head())**

	ID	Store ID	Total Price	Base Price	Units Sold
0	1	8091	99.0375	111.8625	20
1	2	8091	99.0375	99.0375	28
2	3	8091	133.9500	133.9500	19
3	4	8091	133.9500	133.9500	44
4	5	8091	141.0750	141.0750	52

**print(df.tail())**

	ID	Store ID	Total Price	Base Price	Units Sold
150145	212638	9984	235.8375	235.8375	38
150146	212639	9984	235.8375	235.8375	30
150147	212642	9984	357.6750	483.7875	31
150148	212643	9984	141.7875	191.6625	12
150149	212644	9984	234.4125	234.4125	15

**print(df.info())**

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150150 entries, 0 to 150149
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   ID           150150 non-null  int64
1   Store ID     150150 non-null  int64
2   Total Price  150149 non-null  float64
3   Base Price   150150 non-null  float64
4   Units Sold   150150 non-null  int64
dtypes: float64(2), int64(3)
memory usage: 5.7 MB
None
```

```
print(df.describe())
```

```
ID      Store ID  Total Price  Base Price \
count  150150.000000  150150.000000  150149.000000  150150.000000
mean    106271.555504   9199.422511   206.626751   219.425927
std     61386.037861   615.591445    103.308516   110.961712
min       1.000000    8023.000000    41.325000    61.275000
25%     53111.250000   8562.000000   130.387500   133.237500
50%    106226.500000   9371.000000   198.075000   205.912500
75%    159452.750000   9731.000000   233.700000   234.412500
max    212644.000000  9984.000000   562.162500   562.162500
```

```
Units Sold
count  150150.000000
mean     51.674206
std     60.207904
min       1.000000
25%     20.000000
50%     35.000000
75%     62.000000
max    2876.000000
```

## Identifying null Values

```
print(df.isnull())
```

```
ID  Store ID  Total Price  Base Price  Units Sold
0   False    False      False      False      False
1   False    False      False      False      False
2   False    False      False      False      False
3   False    False      False      False      False
4   False    False      False      False      False
...   ...     ...        ...        ...        ...
150145 False    False      False      False      False
150146 False    False      False      False      False
150147 False    False      False      False      False
150148 False    False      False      False      False
150149 False    False      False      False      False
```

```
[150150 rows x 5 columns]
```

```
c = df.isnull().sum()
print(c)
```

```
ID      0
Store ID  0
Total Price  1
Base Price  0
Units Sold  0
dtype: int64
```

```
print('Total Sum of null values in the Data set = ',c.sum())
Total Sum of null values in the Data set = 1
```

### Data Preprocessing - Replacing the null values

```
df.drop_duplicates()
```

ID	Store ID	Total Price	Base Price	Units Sold	
0	1	8091	99.0375	111.8625	20
1	2	8091	99.0375	99.0375	28
2	3	8091	133.9500	133.9500	19
3	4	8091	133.9500	133.9500	44
4	5	8091	141.0750	141.0750	52
...	...	...	...	...	...
150145	212638	9984	235.8375	235.8375	38
150146	212639	9984	235.8375	235.8375	30
150147	212642	9984	357.6750	483.7875	31
150148	212643	9984	141.7875	191.6625	12
150149	212644	9984	234.4125	234.4125	15

150150 rows × 5 columns

**df.fillna(0)**

ID	Store ID	Total Price	Base Price	Units Sold	
0	1	8091	99.0375	111.8625	20
1	2	8091	99.0375	99.0375	28
2	3	8091	133.9500	133.9500	19
3	4	8091	133.9500	133.9500	44
4	5	8091	141.0750	141.0750	52
...	...	...	...	...	...
150145	212638	9984	235.8375	235.8375	38
150146	212639	9984	235.8375	235.8375	30
150147	212642	9984	357.6750	483.7875	31
150148	212643	9984	141.7875	191.6625	12
150149	212644	9984	234.4125	234.4125	15

150150 rows × 5 columns

## Data Normalization

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
df['Values_standardized'] = scaler.fit_transform(df[['Total Price']])
df['Values_standardized1'] = scaler.fit_transform(df[['Base Price']])
print(df)
```

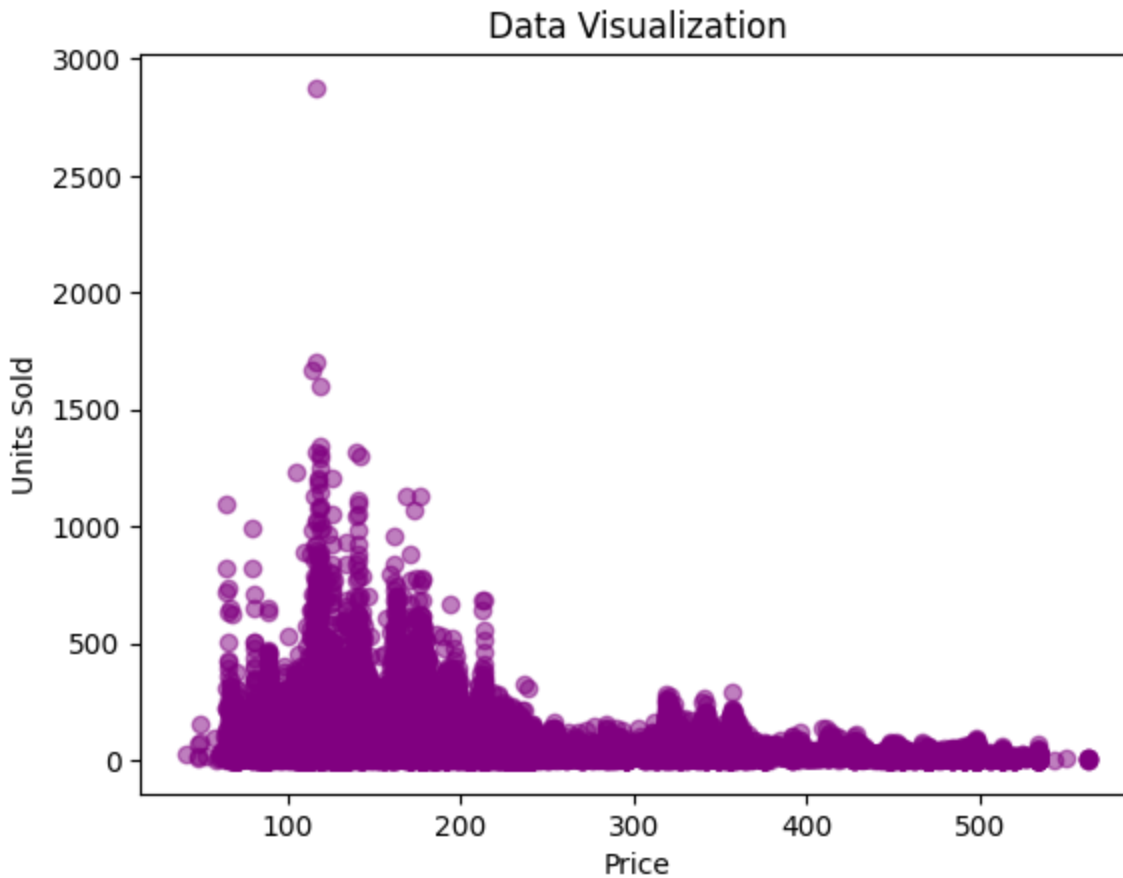
ID	Store ID	Total Price	Base Price	Units Sold	\
0	1	8091	99.0375	111.8625	20
1	2	8091	99.0375	99.0375	28
2	3	8091	133.9500	133.9500	19
3	4	8091	133.9500	133.9500	44
4	5	8091	141.0750	141.0750	52
...	...	...	...	...	...
150145	212638	9984	235.8375	235.8375	38
150146	212639	9984	235.8375	235.8375	30
150147	212642	9984	357.6750	483.7875	31
150148	212643	9984	141.7875	191.6625	12
150149	212644	9984	234.4125	234.4125	15

	Values_standardized	Values_standardized1
0	-1.041440	-0.969377
1	-1.041440	-1.084958
2	-0.703495	-0.770322
3	-0.703495	-0.770322
4	-0.634526	-0.706110
...	...	...
150145	0.282754	0.147904
150146	0.282754	0.147904
150147	1.462113	2.382466
150148	-0.627629	-0.250208
150149	0.268960	0.135061

[150150 rows x 7 columns]

## Data Visualisation

```
plt.scatter(df['Total Price'], df['Units Sold'], alpha=0.5, color='purple')
plt.title('Data Visualization')
plt.xlabel('Price')
plt.ylabel('Units Sold')
plt.show()
```



## Data Analysis using different models

### Simple Linear Regression

```
cdf = df[['Total Price','Base Price','Units Sold']]
```

```
cdf.head(9)
```

	Total Price	Base Price	Units Sold
0	99.0375	111.8625	20
1	99.0375	99.0375	28
2	133.9500	133.9500	19
3	133.9500	133.9500	44
4	141.0750	141.0750	52
5	227.2875	227.2875	18
6	327.0375	327.0375	47
7	210.9000	210.9000	50
8	190.2375	234.4125	82



```

test = test.dropna()
y_hat= regr.predict(test[['Total Price', 'Base Price']])
x = np.asanyarray(test[['Total Price', 'Base Price']])
y = np.asanyarray(test[['Units Sold']])
print("Mean Squared Error (MSE) : %.2f"
% np.mean((y_hat - y) ** 2))
# Explained variance score: 1 is perfect prediction
print('Variance score: %.2f' % regr.score(x, y))

```

Mean Squared Error (MSE) : 2948.93

Variance score: 0.15

```

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:432: UserWarning: X has
feature names, but LinearRegression was fitted without feature names
warnings.warn(

```

## Random forest Algorithm

```


import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt
# Load your dataset from the CSV file
data = pd.read_csv("PoductDemand.csv") # Adjust the
# Handle missing values by filling with the mean
data = data.fillna(data.mean())
features = data[['Total Price', 'Base Price']]
target = data['Units Sold']
# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(features, target,
test_size=0.2, random_state=42)
# Create and train the Random Forest model
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)
# Make predictions on the test set
y_pred = rf_model.predict(X_test)
# Calculate model performance metrics
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

```

```
print(f"Mean Squared Error: {mse:.2f}")
print(f"R-squared (R2) Score: {r2:.2f}")
# Visualize the feature importances
feature_importances = rf_model.feature_importances_
plt.bar(features.columns, feature_importances)
plt.xlabel("Features")
plt.ylabel("Importance")
plt.title("Feature Importances")
plt.show()
```

Mean Squared Error: 1885.60

R-squared (R2) Score: 0.43




## **Regression with the Gradient Boosting algorithm**

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt
# Load your dataset from the CSV file
data = pd.read_csv("PoductDemand.csv")
# Handle missing values by filling with the mean
data = data.fillna(data.mean())
# Select the relevant features (SO2 and NO2) and the target variable
features = data[['Total Price', 'Base Price']]
target = data['Units Sold']
# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(features, target,
test_size=0.2, random_state=42)
# Create and train the Gradient Boosting model
gb_model = GradientBoostingRegressor(n_estimators=100,
random_state=42)
gb_model.fit(X_train, y_train)
# Make predictions on the test set
y_pred = gb_model.predict(X_test)
# Calculate model performance metrics
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f"Mean Squared Error: {mse:.2f}")
print(f"R-squared (R2) Score: {r2:.2f}")
# Visualize the feature importances
feature_importances = gb_model.feature_importances_
plt.bar(features.columns, feature_importances)
plt.xlabel("Features")
plt.ylabel("Importance")
plt.title("Feature Importances")
plt.show()
```

Mean Squared Error: 1987.34

R-squared (R<sup>2</sup>) Score: 0.39



## Model Tuning

Fine-tuning of the selected model(s) will be carried out to optimize their hyperparameters and enhance their predictive power.

## Model Deployment

Once a satisfactory model is developed, we will deploy it to make real-time predictions on new data, allowing the business to anticipate product demand more effectively.

## **Conclusion**

This data science project aims to provide a solution for product demand prediction, which can have significant implications for inventory management and sales optimization. The final outcome of the project will be a machine learning model capable of forecasting product demand based on the given dataset.